

Analyzing the NYC Subway Dataset

Intro to Data Science: Final Project 1, Part 2

(Short Questions)

Data Exploration Supplement

Austin J. Alexander

Import Directives and Initial DataFrame Creation

```
In [40]: import inflect # for string manipulation
import numpy as np
import pandas as pd
import scipy as sp
import scipy.stats as st
import matplotlib.pyplot as plt
%matplotlib inline

filename = '/Users/excalibur/py/nanodegree/intro_ds/final_project/improved-dataset/turnstile_weather_v2.csv'

# import data
data = pd.read_csv(filename)
```

Initial Data Exploration

Data Shape

```
In [41]: print "SHAPE: " + str(data.shape)
data.head(1)
```

SHAPE: (42649, 27)

Out[41]:

	UNIT	DATE _n	TIME _n	ENTRIES _n	EXITS _n	ENTRIES _n _hourly	EXITS _n _hourly
0	R003	05-01-11	00:00:00	4388333	2911002	0	0

1 rows × 27 columns

Data Types

```
In [42]: data.dtypes
```

```
Out[42]: UNIT                object
DATEn                object
TIMEn                object
ENTRIESn             int64
EXITSn               int64
ENTRIESn_hourly       float64
EXITSn_hourly         float64
datetime                   object
hour                       int64
day_week                   int64
weekday                   int64
station                    object
latitude                   float64
longitude                  float64
conds                      object
fog                        int64
precipi                    float64
pressurei                  float64
rain                       int64
tempi                      float64
wspdi                      float64
meanprecipi                float64
meanpressurei              float64
meantempi                 float64
meanwspdi                  float64
weather_lat                float64
weather_lon                float64
dtype: object
```

ENTRIESn_hourly Statistics

```
In [43]: data['ENTRIESn_hourly'].describe()
```

```
Out[43]: count      42649.000000  
         mean        1886.589955  
         std         2952.385585  
         min          0.000000  
         25%         274.000000  
         50%         905.000000  
         75%        2255.000000  
         max        32814.000000  
         Name: ENTRIESn_hourly, dtype: float64
```

Functions for Getting, Mapping, and Plotting Data

[N.B. Due to decisions described in the [*Unit Entries* supplement \(IntroDS-ProjectOne-Unit Entries-Supplement.ipynb\)](#), in the current analysis, unless otherwise noted, *entries* will refer to a summation of **ENTRIESn_hourly** per **UNIT** (i.e., not, as might be expected, simply alues in the **ENTRIESn** column).]

```
In [44]: entries_hourly_by_row = data['ENTRIESn_hourly'].values
```

```
In [45]: def map_column_to_entries_hourly(column):  
         instances = column.values # e.g., longitude_instances = data['lon  
         gitude'].values  
  
         # reduce  
         entries_hourly = {} # e.g., longitude_entries_hourly = {}  
         for i in np.arange(len(instances)):  
             if instances[i] in entries_hourly:  
                 entries_hourly[instances[i]] += float(entries_hourly_b  
y_row[i])  
             else:  
                 entries_hourly[instances[i]] = float(entries_hourly_b  
y_row[i])  
  
         return entries_hourly # e.g., longitudes, entries
```

```
In [46]: def display_basic_stats(entries_hourly_dict, column1name):
    # e.g, longitude_df = pd.DataFrame(data=longitude_entries_hourly.items(), columns=['longitude', 'entries'])
    df = pd.DataFrame(data=entries_hourly_dict.items(), columns=[column1name, 'entries'])

    p = inflect.engine()
    print "{0} AND THEIR ENTRIES".format(p.plural(column1name.upper()))
    print df.head(3)

    print
    print pd.DataFrame(df['entries']).describe()
    print "{:<7}".format('range') + "{:0<14}".format(str(np.ptp(entries_hourly_dict.values()))))

    return df # e.g, longitude_df
```

```
In [47]: def plot_data(df, column1name, plot_kind, xaxis_labeled):

    p = inflect.engine()
    if xaxis_labeled == True:
        df.plot(x=column1name, y='entries', title="{0} AND THEIR ENTRIES".format(p.plural(column1name.upper()))), kind=plot_kind, alpha=0.5, color='green')
        plt.xlabel(column1name)
    else:
        df.plot(title="{0} AND THEIR ENTRIES".format(p.plural(column1name.upper()))), kind=plot_kind, alpha=0.5, color='green')
        plt.xlabel("{0} row index".format(column1name))

    plt.ylabel('{0} entries'.format(column1name))
    plt.legend(['entries'])
    plt.show()
```

```
In [48]: def plot_histogram(df, column_name, num_of_bins):
    df[column_name].plot(kind='hist', bins=num_of_bins, alpha=0.5, color='green')
    plt.ylabel('frequency')
    plt.show()
```

Function for Basic Statistics

```
In [49]: def describe_samples(sample1, sample2):
    size1, min_max1, mean1, var1, skew1, kurt1 = st.describe(sample
1)
    size2, min_max2, mean2, var2, skew2, kurt2 = st.describe(sample
2)

    med1 = np.median(sample1)
    med2 = np.median(sample2)

    std1 = np.std(sample1)
    std2 = np.std(sample2)

    print "Sample 1 (rainy days):\n min = {0}, max = {1},\n mean
= {2:.2f}, median = {3}, var = {4:.2f}, std = {5:.2f}".format(min_m
ax1[0], min_max1[1], mean1, med1, var1, std1)
    print "Sample 2 (non-rainy days):\n min = {0}, max = {1},\n m
ean = {2:.2f}, median = {3}, var = {4:.2f}, std = {5:.2f}".format(m
in_max2[0], min_max2[1], mean2, med2, var2, std2)
```

Non-Weather-Related Data

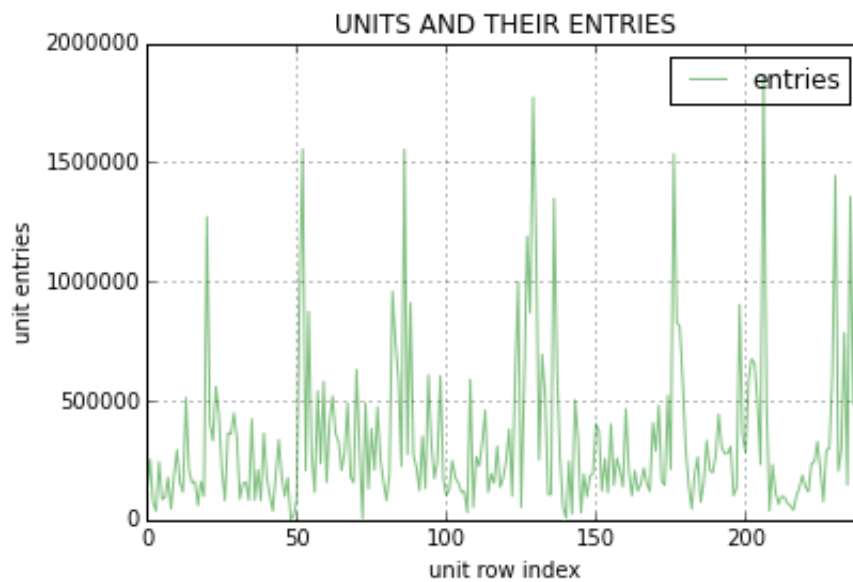
Unit Statistics

```
In [50]: unit_entries_hourly = map_column_to_entries_hourly(data['UNIT'])
unit_df = display_basic_stats(unit_entries_hourly, 'unit')
plot_data(unit_df, 'unit', 'line', False)
```

UNITS AND THEIR ENTRIES

	unit	entries
0	R318	112098
1	R319	254531
2	R312	73913

	entries
count	240.000000
mean	335254.895833
std	334849.388932
min	0.000000
25%	131148.000000
50%	221479.500000
75%	409285.750000
max	1868674.000000
range	1868674.000000



Top-5 Units

```
In [51]: unit_df.sort(columns='entries', ascending=False).head(5)
```

Out[51]:

	unit	entries
206	R084	1868674
129	R022	1773372
237	R012	1618262
52	R046	1555117
86	R055	1554806

Unit Summary

Clearly, certain units received more entries than other units.

Date Statistics

```
In [52]: date_entries_hourly = map_column_to_entries_hourly(data['DATEn'])
date_df = display_basic_stats(date_entries_hourly, 'date')
plot_data(date_df, 'date', 'line', False)
```

DATES AND THEIR ENTRIES

	date	entries
0	05-30-11	1409572
1	05-15-11	1413929
2	05-04-11	3118915

	entries
count	31.000000
mean	2595521.774194
std	710440.834289
min	1400098.000000
25%	1891834.000000
50%	3009536.000000
75%	3137683.000000
max	3201840.000000
range	1801742.000000

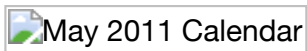
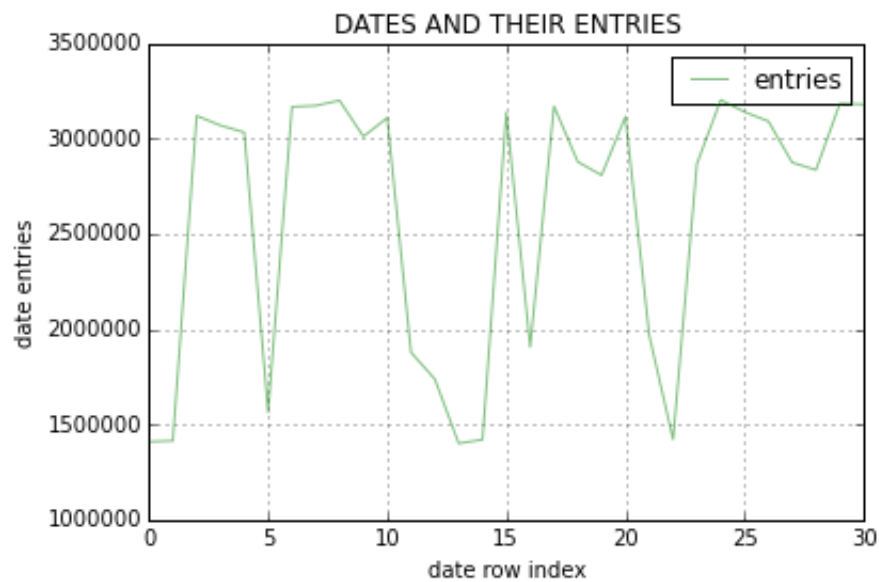


Image from: calendarbar.org

Top-5 Dates


```
In [53]: date_df.sort(columns='entries', ascending=False).head(5)
```

```
Out[53]:
```

	date	entries
24	05-12-11	3201840
8	05-05-11	3199002
29	05-03-11	3183128
30	05-06-11	3179032
7	05-26-11	3172004

Date Summary

Clearly, certain dates received more entries than other dates (the top 5 included 1 Tuesday, 3 Thursdays, and 1 Friday).

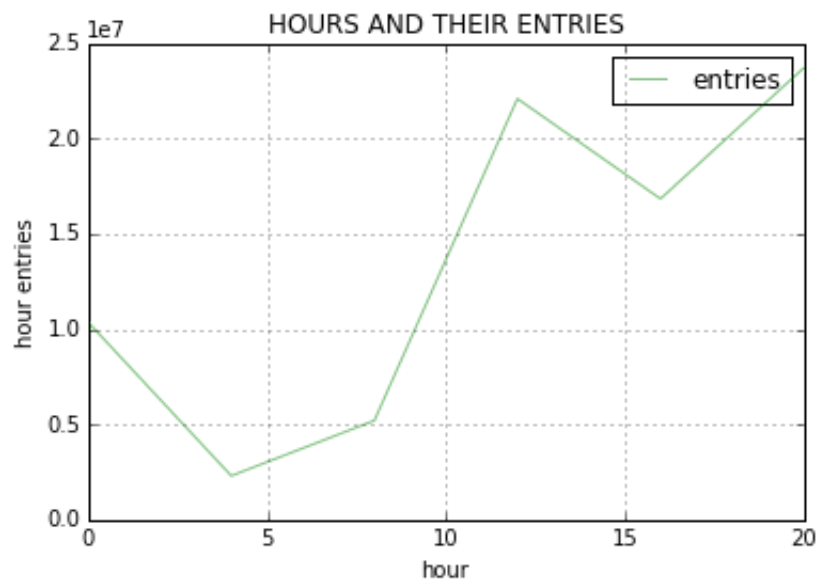
Hour Statistics

```
In [54]: hour_entries_hourly = map_column_to_entries_hourly(data['hour'])
hour_df = display_basic_stats(hour_entries_hourly, 'hour')
plot_data(hour_df, 'hour', 'line', True)
```

HOURS AND THEIR ENTRIES

	hour	entries
0	0	10353167
1	4	2300788
2	8	5198583

	entries
count	6.000000
mean	13410195.833333
std	8863957.086415
min	2300788.000000
25%	6487229.000000
50%	13593103.500000
75%	20772247.000000
max	23690281.000000
range	21389493.000000



Top-5 Hours

```
In [55]: hour_df.sort(columns='entries', ascending=False).head(5)
```

Out[55]:

	hour	entries
5	20	23690281
3	12	22085316
4	16	16833040
0	0	10353167
2	8	5198583

Hour Summary

Clearly, certain hours recorded more entries than other hours (the top 2 were, in order, 10:00pm and 12:00pm).

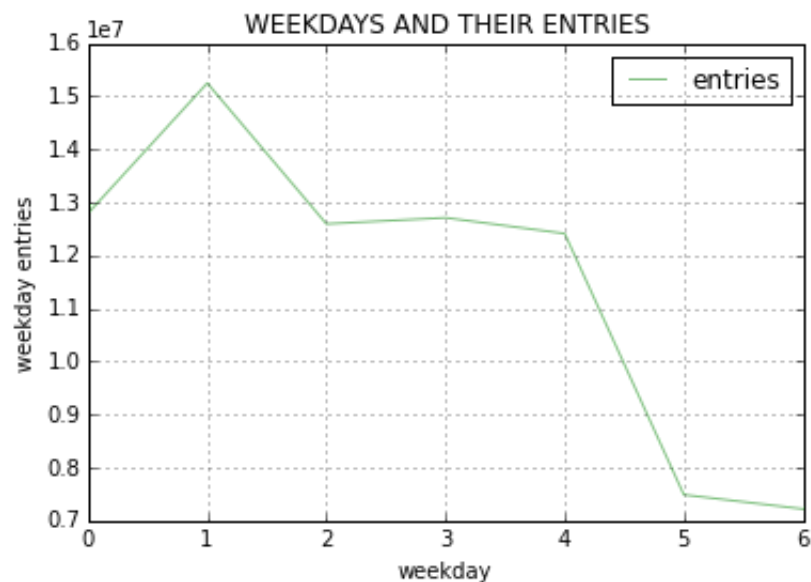
Weekday Statistics

```
In [56]: weekday_entries_hourly = map_column_to_entries_hourly(data['day_week'])
weekday_df = display_basic_stats(weekday_entries_hourly, 'weekday')
plot_data(weekday_df, 'weekday', 'line', True)
```

WEEKDAYS AND THEIR ENTRIES

	weekday	entries
0	0	12795107
1	1	15246943
2	2	12592691

	entries
count	7.000000
mean	11494453.571429
std	2989933.638739
min	7218706.000000
25%	9949293.000000
50%	12592691.000000
75%	12752124.500000
max	15246943.000000
range	8028237.000000



Top-5 Weekdays

(0: Mon, 1: Tue, 2: Wed, 3: Thu, 4: Fri, 5: Sat, 6: Sun)

```
In [57]: weekday_df.sort(columns='entries', ascending=False).head(5)
```

```
Out[57]:
```

	weekday	entries
1	1	15246943
0	0	12795107
3	3	12709142
2	2	12592691
4	4	12411679

Weekday Summary

Clearly, certain weekdays received more entries than other weekdays (the top 2 were, in order, Tuesday and Monday -- somewhat strange given the Date results above).

Station Statistics

[N.B. Some stations have the same name but different locations, so unique identifiers needed to be created.]

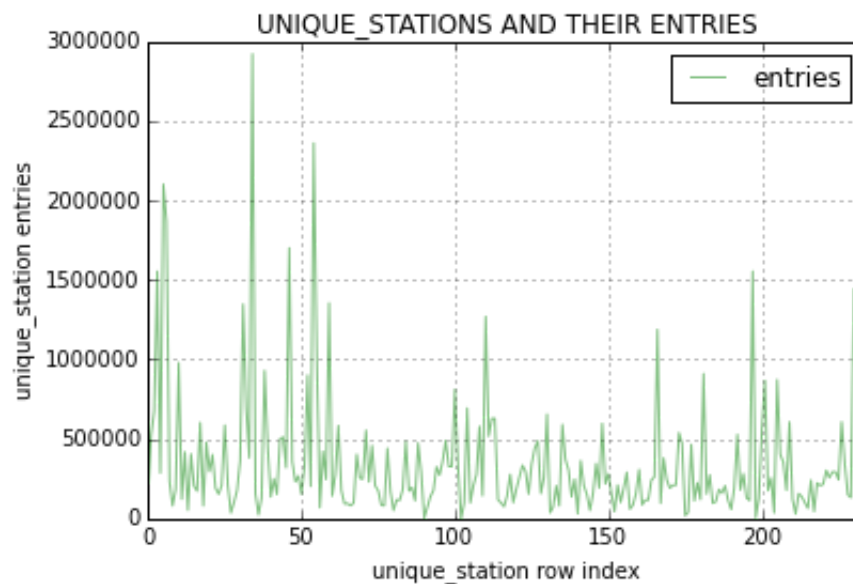
```
In [58]: data['unique_station'] = data['station'] + " (" + data['latitude'].map(str) + ", " + data['longitude'].map(str) + ")"
```

```
In [59]: station_entries_hourly = map_column_to_entries_hourly(data['unique_station'])
station_df = display_basic_stats(station_entries_hourly, 'unique_station')
plot_data(station_df, 'unique_station', 'line', False)
```

UNIQUE_STATIONS AND THEIR ENTRIES

	unique_station	entries
0	176 ST (40.848635, -73.912497)	151399
1	168 ST-BROADWAY (40.840778, -73.940091)	521054
2	57 ST-7 AVE (40.764755, -73.980646)	674799

	entries
count	234.000000
mean	343851.175214
std	393424.158576
min	0.000000
25%	130422.000000
50%	217648.000000
75%	402551.250000
max	2920887.000000
range	2920887.000000



Station Summary

Clearly, certain stations received more entries than other stations.

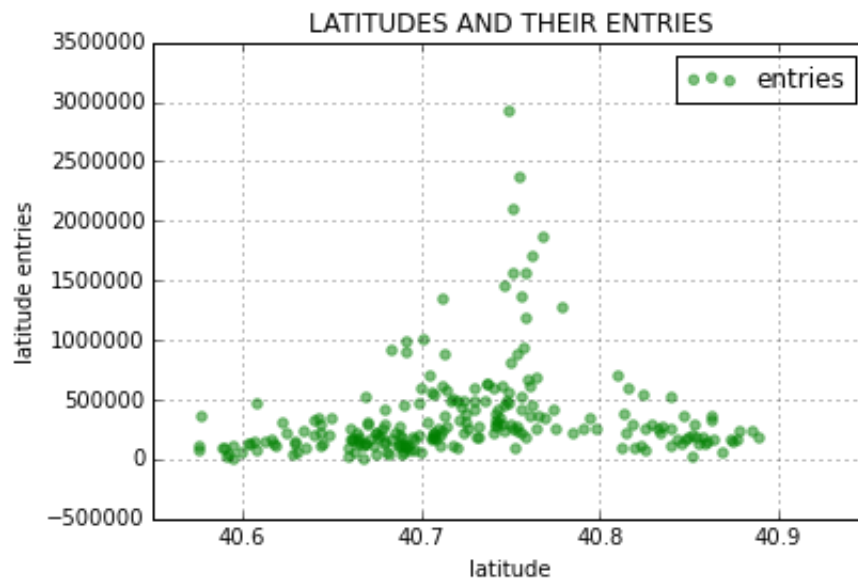
Latitude Statistics

```
In [60]: latitude_entries_hourly = map_column_to_entries_hourly(data['latitude'])
latitude_df = display_basic_stats(latitude_entries_hourly, 'latitude')
plot_data(latitude_df, 'latitude', 'scatter', True)
```

LATITUDES AND THEIR ENTRIES

	latitude	entries
0	40.852417	7559
1	40.707840	209745
2	40.643982	102508

	entries
count	233.000000
mean	345326.931330
std	393653.267874
min	0.000000
25%	131511.000000
50%	218938.000000
75%	402883.000000
max	2920887.000000
range	2920887.000000



Latitude Summary

Clearly, certain latitudes received more entries than other latitudes (in particular, those between latitudes 40.7 and 40.8).

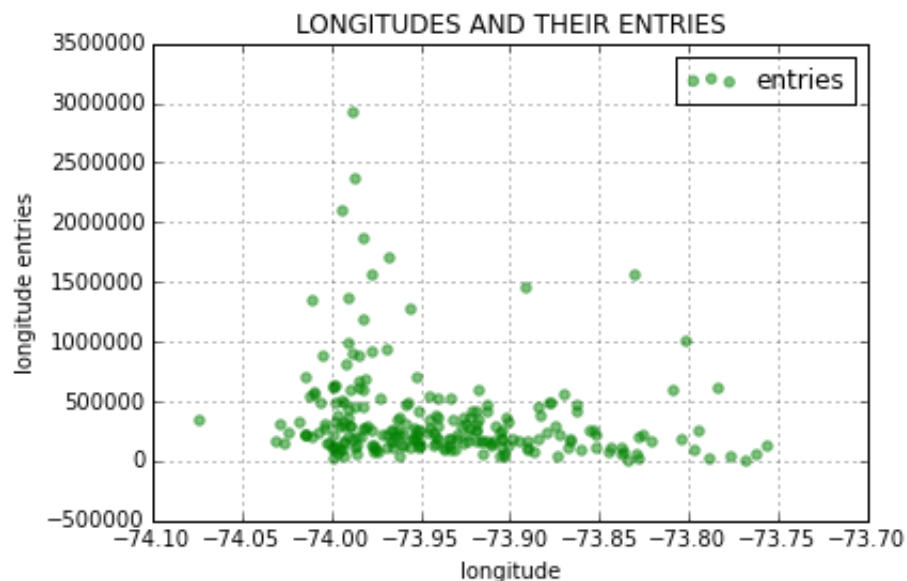
Longitude Statistics

```
In [61]: longitude_entries_hourly = map_column_to_entries_hourly(data['longitude'])
longitude_df = display_basic_stats(longitude_entries_hourly, 'longitude')
plot_data(longitude_df, 'longitude', 'scatter', True)
```

LONGITUDES AND THEIR ENTRIES

	longitude	entries
0	-73.977417	911174
1	-73.828125	193792
2	-74.014099	694605

	entries
count	234.000000
mean	343851.175214
std	393424.158576
min	0.000000
25%	130422.000000
50%	217648.000000
75%	402551.250000
max	2920887.000000
range	2920887.000000



Longitude Summary

Clearly, certain longitudes received more entries than other longitudes (in particular, those between longitudes -74.00 and -73.95).

Combining Station, Latitude, Longitude Data on a Map Layer

The top-ten most-entered stations are the ones that someone familiar with New York City might expect (with a focus on Manhattan Island):

```
In [62]: station_location = data[['unique_station', 'latitude', 'longitude']]
station_location.drop_duplicates(inplace=True)
station_location_entries = pd.merge(station_location, station_df, on='unique_station')
station_location_entries.sort(columns='entries', ascending=False, inplace=True)
top_ten = station_location_entries.head(10)
print top_ten[['unique_station', 'entries']]

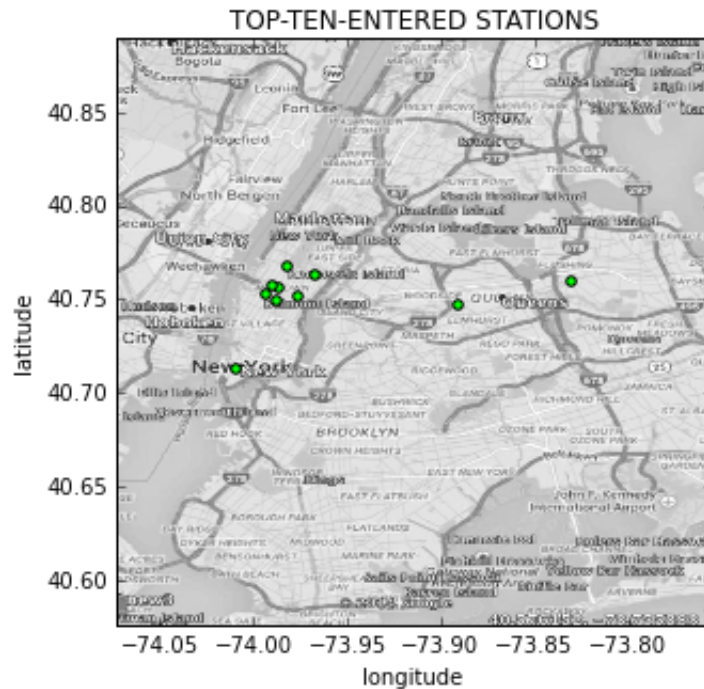
plt.figure(figsize = (5,5))
plt.title('TOP-TEN-ENTERED STATIONS')
plt.xlabel('longitude')
plt.ylabel('latitude')
plt.xlim(station_location_entries['longitude'].min(), station_location_entries['longitude'].max())
plt.ylim(station_location_entries['latitude'].min(), station_location_entries['latitude'].max())
img = plt.imread('NYmap.png')

plt.scatter(top_ten['longitude'], top_ten['latitude'], color='#00FF00', edgecolors='black', zorder=1)

plt.imshow(img, zorder=0, extent=[station_location_entries['longitude'].min(), station_location_entries['longitude'].max(), station_location_entries['latitude'].min(), station_location_entries['latitude'].max()])

plt.show()
```

			unique_station	entries
14	34	ST-HERALD SQ	(40.749533, -73.987899)	2920887
21	42	ST-TIMES SQ	(40.755905, -73.986504)	2360981
8	34	ST-PENN STA	(40.752247, -73.993456)	2101634
58	59	ST-COLUMBUS	(40.76811, -73.981891)	1868674
34		LEXINGTON AVE	(40.762796, -73.967686)	1701440
32	42	ST-GRD CNTRL	(40.751849, -73.976945)	1555117
38		MAIN ST	(40.759578, -73.830056)	1554806
10		ROOSEVELT AVE	(40.746655, -73.891361)	1444569
7	42	ST-PA BUS TE	(40.757303, -73.989787)	1355492
18		WORLD TRADE CTR	(40.712557, -74.009807)	1347727



Station, Latitude, and Longitude Summary

Again, it seems quite clear that number of entries is dependent on *location, location, location!*

Weather-Related Data

While the above simple observations seem to indicate that non-weather-related variables have a tremendous impact on transit usage, it would only be fair to explore possible weather-related influences as well (esp. since the project guidelines expect it). Out of the weather-related data that has been made available, the two most-sensible categories to check are the (binary) rain and temperature (in Fahrenheit) variables.

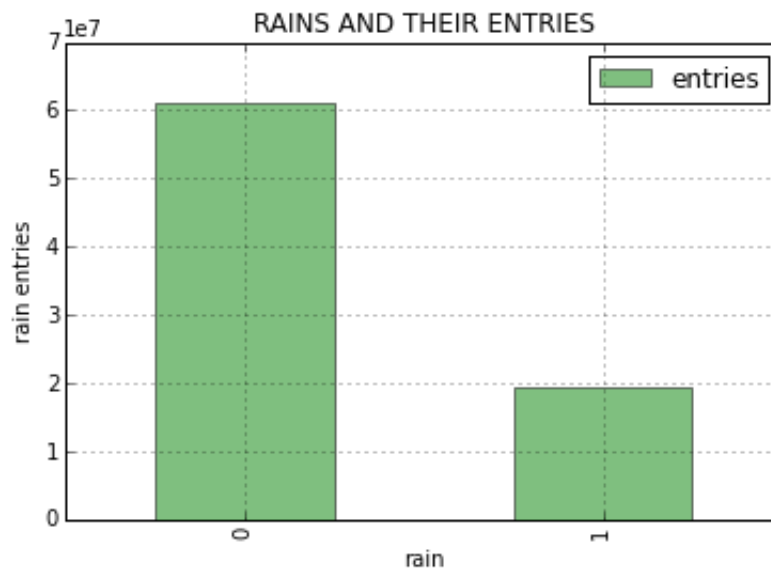
Rain Statistics

```
In [63]: rain_entries_hourly = map_column_to_entries_hourly(data['rain'])
rain_df = display_basic_stats(rain_entries_hourly, 'rain')
plot_data(rain_df, 'rain', 'bar', True)
```

RAINS AND THEIR ENTRIES

	rain	entries
0	0	61020916
1	1	19440259

	entries
count	2.000000
mean	40230587.500000
std	29401964.530892
min	19440259.000000
25%	29835423.250000
50%	40230587.500000
75%	50625751.750000
max	61020916.000000
range	41580657.000000



```
In [64]: rain_days = data[data['rain'] == 1]
no_rain_days = data[data['rain'] == 0]

print "RAIN DAYS"
print rain_days['ENTRIESn_hourly'].describe()
print
print "NO-RAIN DAYS"
print no_rain_days['ENTRIESn_hourly'].describe()
```

```
RAIN DAYS
count      9585.000000
mean       2028.196035
std        3189.433373
min         0.000000
25%        295.000000
50%        939.000000
75%       2424.000000
max       32289.000000
Name: ENTRIESn_hourly, dtype: float64
```

```
NO-RAIN DAYS
count      33064.000000
mean       1845.539439
std        2878.770848
min         0.000000
25%        269.000000
50%        893.000000
75%       2197.000000
max       32814.000000
Name: ENTRIESn_hourly, dtype: float64
```

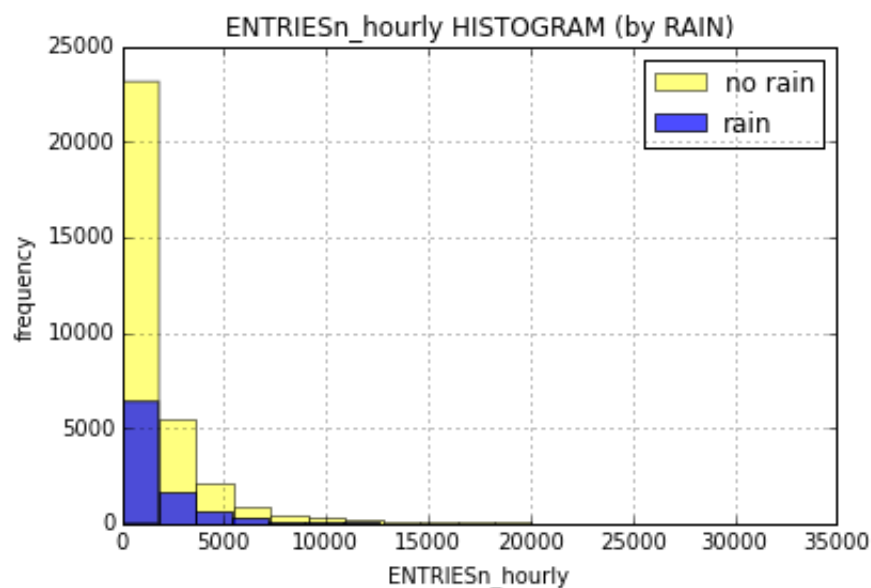
ENTRIESn_hourly HISTOGRAM (by RAIN)

```
In [76]: print "Number of non-rainy days:", no_rain_days['ENTRIESn_hourly'].count()
print "Number of rainy days:", rain_days['ENTRIESn_hourly'].count()

no_rain_days['ENTRIESn_hourly'].plot(kind='hist', bins=18, alpha=0.5, color='yellow')
rain_days['ENTRIESn_hourly'].plot(kind='hist', bins=18, alpha=0.7, color='blue')
plt.title('ENTRIESn_hourly HISTOGRAM (by RAIN)')
plt.xlabel('ENTRIESn_hourly')
plt.ylabel('frequency')
plt.legend(['no rain', 'rain'])
plt.show()
```

Number of non-rainy days: 33064

Number of rainy days: 9585



```

In [66]: date_and_rain = data[['DATEn', 'rain']].drop_duplicates()
date_and_rain.sort(columns='DATEn', inplace=True)
print date_and_rain.head()

dates = data['DATEn'].unique()
rain_dates = date_and_rain[date_and_rain['rain'] == 1]['DATEn'].unique()
no_rain_dates = date_and_rain[date_and_rain['rain'] == 0]['DATEn'].unique()

indices_of_rain_dates = []
for rain_date in rain_dates:
    indices_of_rain_dates.append(np.where(dates == rain_date)[0][0])

indices_of_no_rain_dates = []
for no_rain_date in no_rain_dates:
    indices_of_no_rain_dates.append(np.where(dates == no_rain_date)[0][0])

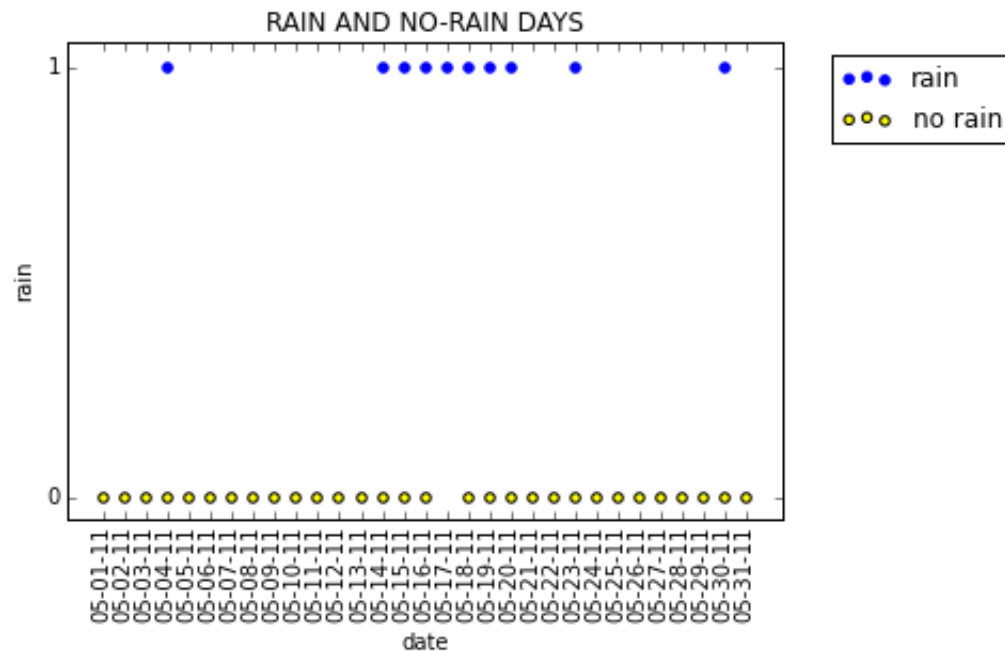
plt.title('RAIN AND NO-RAIN DAYS')
plt.xticks(np.arange(len(dates)), dates, rotation='vertical')
plt.yticks([0,1])
plt.xlabel('date')
plt.ylabel('rain')

plt.scatter(indices_of_rain_dates, np.ones(len(indices_of_rain_dates)), color='blue')
plt.scatter(indices_of_no_rain_dates, np.zeros(len(indices_of_no_rain_dates)), color='yellow', edgecolors='black')

plt.legend(['rain', 'no rain'], bbox_to_anchor=(1.05, 1), loc=2)
plt.show()

```

	DATE	rain
0	05-01-11	0
5	05-02-11	0
11	05-03-11	0
16	05-04-11	1
32542	05-04-11	0



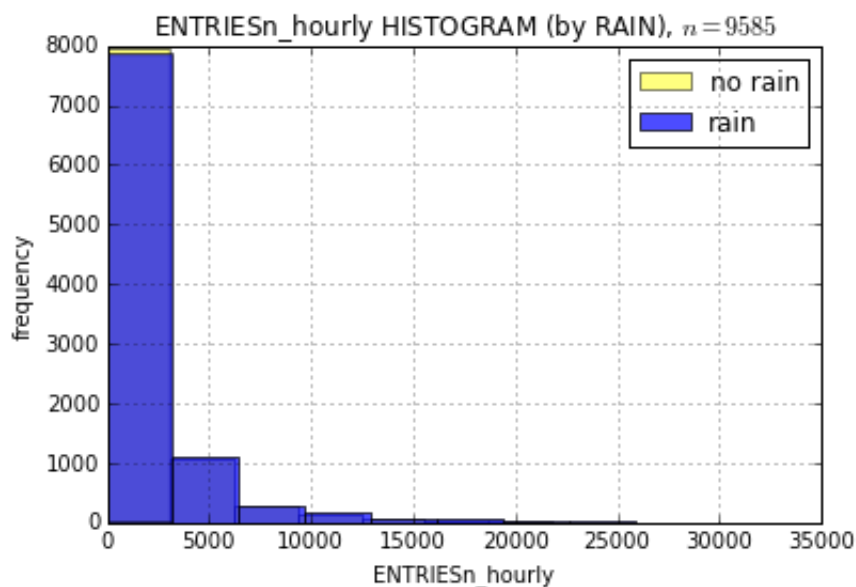
One possible explanation for the difference might be, simply, that there were more non-rainy days in May 2011, as indicated in the graph above. [N.B. Certain days are reported as being rainy and non-rainy days. For a brief exploration of this phenomenon, see the [Rain supplement \(IntroDS-ProjectOne-Rain-Supplement.ipynb\)](#).]

What if the number of non-rainy days was limited to the total number of rainy days in the data set?

```
In [67]: random_row_indices = np.random.choice(no_rain_days['ENTRIESn_hourly'].index.values, 9585)
```



```
In [68]: no_rain_days['ENTRIESn_hourly'].loc[random_row_indices].plot(kind='hist', bins=10, alpha=0.5, color='yellow')
rain_days['ENTRIESn_hourly'].plot(kind='hist', bins=10, alpha=0.7, color='blue')
plt.title(r'ENTRIESn_hourly HISTOGRAM (by RAIN), $n = 9585$')
plt.xlabel('ENTRIESn_hourly')
plt.ylabel('frequency')
plt.legend(['no rain', 'rain'])
plt.show()
```



With the number of samples and bins equal, the similarities between the two groups are obvious.

Rain Summary

While non-rainy days occur in greater number in this data set (thus, contributing to their higher-frequency counts), the distribution of *ENTRIESn_hourly* for rainy and non-rainy days seems otherwise comparable according to the above histograms.

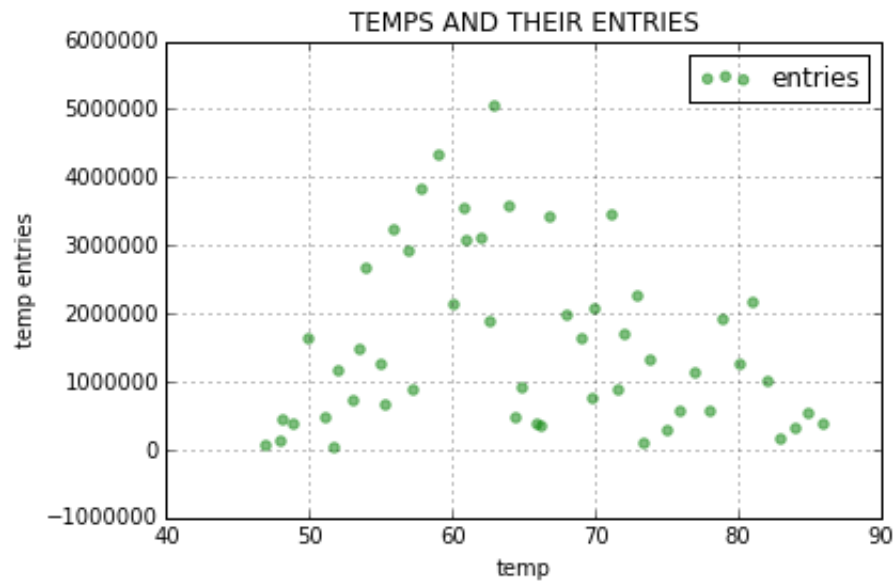
Temperature Statistics

```
In [69]: temp_entries_hourly = map_column_to_entries_hourly(data['temp'])
temp_df = display_basic_stats(temp_entries_hourly, 'temp')
plot_data(temp_df, 'temp', 'scatter', True)
```

TEMPS AND THEIR ENTRIES

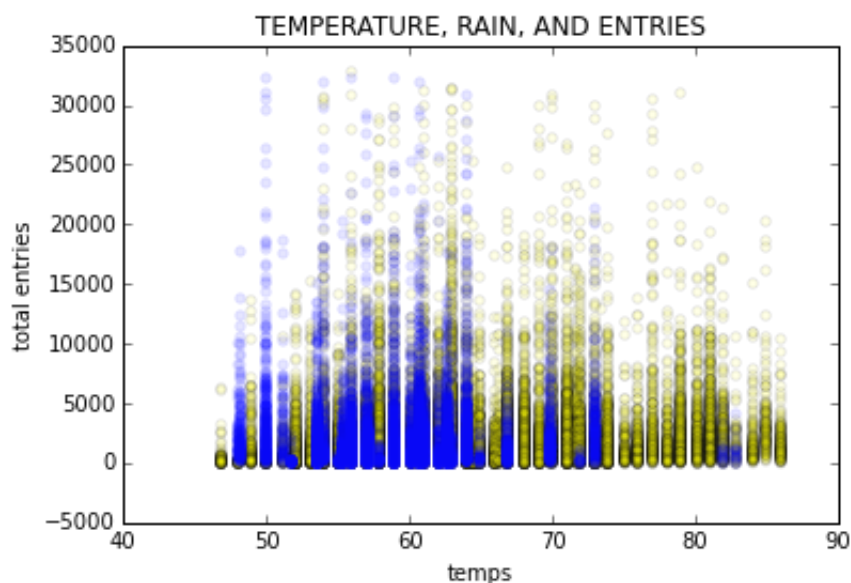
	temp	entries
0	51.1	471715
1	57.9	3831355
2	86.0	362824

	entries
count	52.000000
mean	1547330.288462
std	1276860.663369
min	32822.000000
25%	471789.250000
50%	1203004.000000
75%	2193937.250000
max	5037972.000000
range	5005150.000000



```
In [70]: entries_rain_temp = data[['ENTRIESn_hourly', 'rain', 'temp']].drop_duplicates()
entries_no_rain_temp = entries_rain_temp[entries_rain_temp['rain'] == 0]
entries_rain_temp = entries_rain_temp[entries_rain_temp['rain'] == 1]

plt.scatter(entries_no_rain_temp['temp'], entries_no_rain_temp['ENTRIESn_hourly'], color='yellow', edgecolors='black', alpha=0.1)
plt.scatter(entries_rain_temp['temp'], entries_rain_temp['ENTRIESn_hourly'], color='blue', alpha=0.1)
plt.title('TEMPERATURE, RAIN, AND ENTRIES')
plt.xlabel('temps')
plt.ylabel('total entries')
plt.show()
```



Temperature Summary

Based on the above scatter plots, there does seem to be some type of relationship between temperatures and number of entries. In general, temperatures between 55°F and 66°F received more entries.

Visually-speaking, cold and rainy days seemed to attract approximately the same number of entries as warm and non-rainy (esp. if it's remembered that there were more non-rainy days in the data set).

Preparation for Statistical Tests: Looking for Normality

Taking simple random samples, where n : sample size, and $n = 1000$, the following basic statistics reveal themselves.

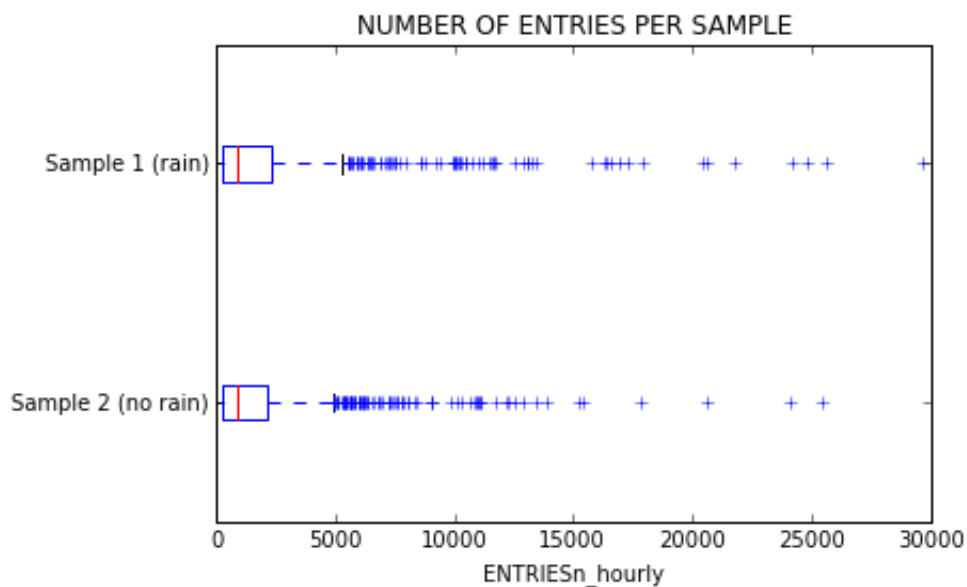
```
In [71]: n = 1000
sample1 = np.random.choice(rain_days['ENTRIESn_hourly'], size=n, replace=False)
sample2 = np.random.choice(no_rain_days['ENTRIESn_hourly'], size=n, replace=False)

describe_samples(sample1, sample2)
```

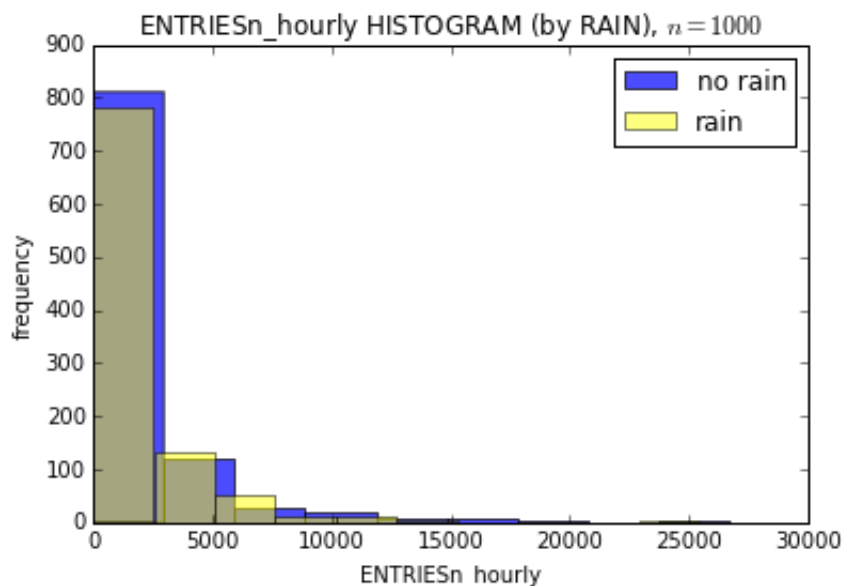
```
Sample 1 (rainy days):
  min = 0.0, max = 29665.0,
  mean = 1975.33, median = 923.0, var = 10059424.13, std = 3170.07
Sample 2 (non-rainy days):
  min = 0.0, max = 25457.0,
  mean = 1803.53, median = 875.5, var = 6855341.46, std = 2616.96
```

```
In [72]: plt.boxplot([sample2, sample1], vert=False)
plt.title('NUMBER OF ENTRIES PER SAMPLE')
plt.xlabel('ENTRIESn_hourly')
plt.yticks([1, 2], ['Sample 2 (no rain)', 'Sample 1 (rain)'])

plt.show()
```



```
In [77]: plt.hist(sample1, color='blue', alpha=0.7, bins=10)
plt.hist(sample2, color='yellow', alpha=0.5, bins=10)
plt.title(r'ENTRIESn_hourly HISTOGRAM (by RAIN), $n = 1000$')
plt.xlabel('ENTRIESn_hourly')
plt.ylabel('frequency')
plt.legend(['no rain', 'rain'])
plt.show()
```



Testing for Normality

Treating the rain and non-rain days as two independent populations, and although visually apparent from the above histogram, the following statistical test seeks to determine whether rainy and non-rainy day distributions are normal.

The Shapiro-Wilk normality test is a test of the null hypothesis that a sample is from a population with a normal distribution.

The test confirms the visually apparent non-normality with a small-enough sample size; so here, $n = 30$.

A 95% level of confidence would suggest that 95% of samples would produce similar statistical results.

For a 95% level of confidence, the level of significance (i.e., the probability of making a Type I error) $\alpha = (1 - 0.05) \cdot 100\% = 0.05$.

```
In [74]: n = 30
alpha = 0.05
small_sample1 = np.random.choice(sample1, size=n, replace=False)
small_sample2 = np.random.choice(sample1, size=n, replace=False)
W1, p1 = st.shapiro(small_sample1)
W2, p2 = st.shapiro(small_sample2)

print "p1 < {0}: {1}".format(alpha, (p1 < alpha))
print "p2 < {0}: {1}".format(alpha, (p2 < alpha))

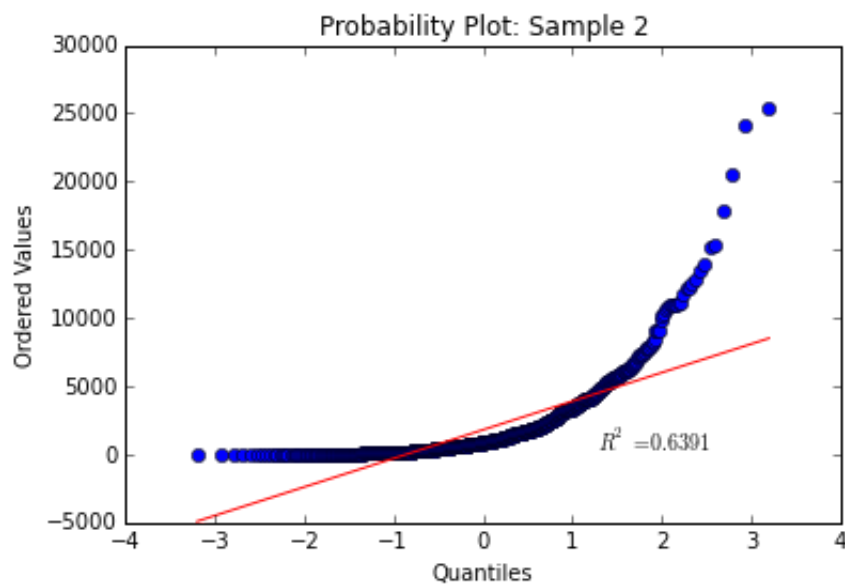
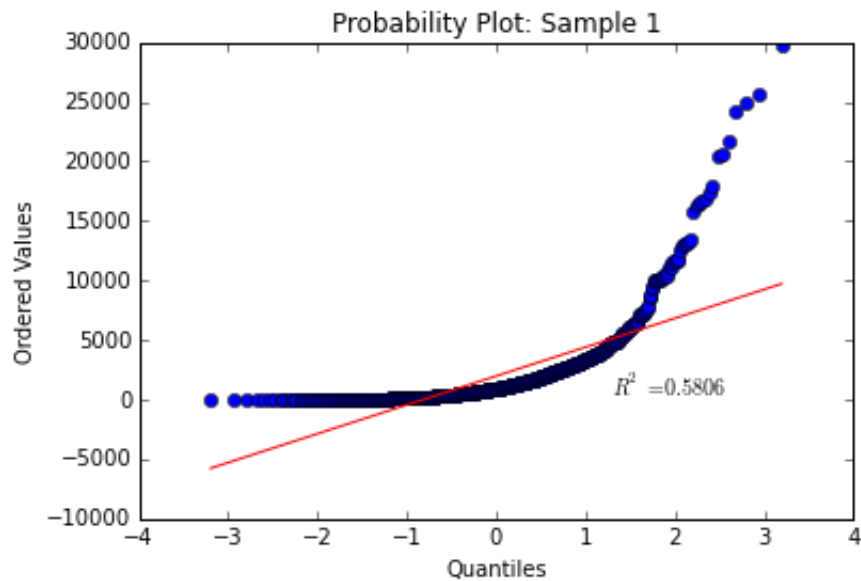
p1 < 0.05: True
p2 < 0.05: True
```

In both cases, $p < 0.05$, so the null hypotheses that the samples come from normally distributed populations are rejected.

Moreover, the following probability plots seal the deal (samples from normal distributions would hug the red regression line throughout the plot):

```
In [75]: st.probplot(sample1, plot=plt)
plt.title('Probability Plot: Sample 1')
plt.show()

st.probplot(sample2, plot=plt)
plt.title('Probability Plot: Sample 2')
plt.show()
```



Apparent Conclusions

Rainy/Non-rainy days and their number of entries are not normally distributed.